


A multidimensional array.

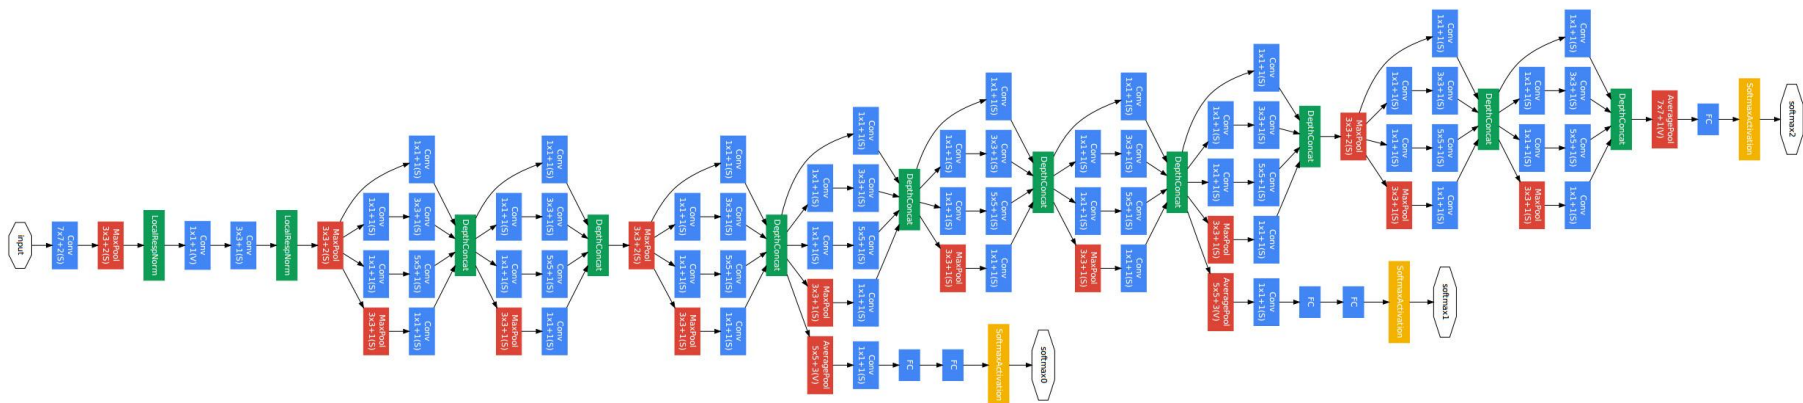


TensorFlow



A graph of operations.

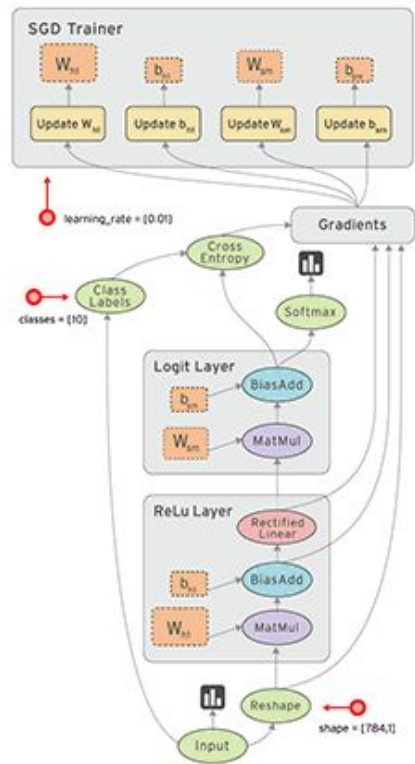
Modeling complexity



The TensorFlow Graph

Computation is defined as a graph

- Graph is defined in high-level language (Python)
- Graph is compiled and optimized
- Graph is executed (in parts or fully) on available low level devices (CPU, GPU, TPU)
- Nodes represent computations and state
- Data (tensors) flow along edges



Build a graph; then run it.


...

```
c = tf.add(a, b)
```


...

```
session = tf.Session()
```

```
value_of_c = session.run(c, {a=1, b=2})
```

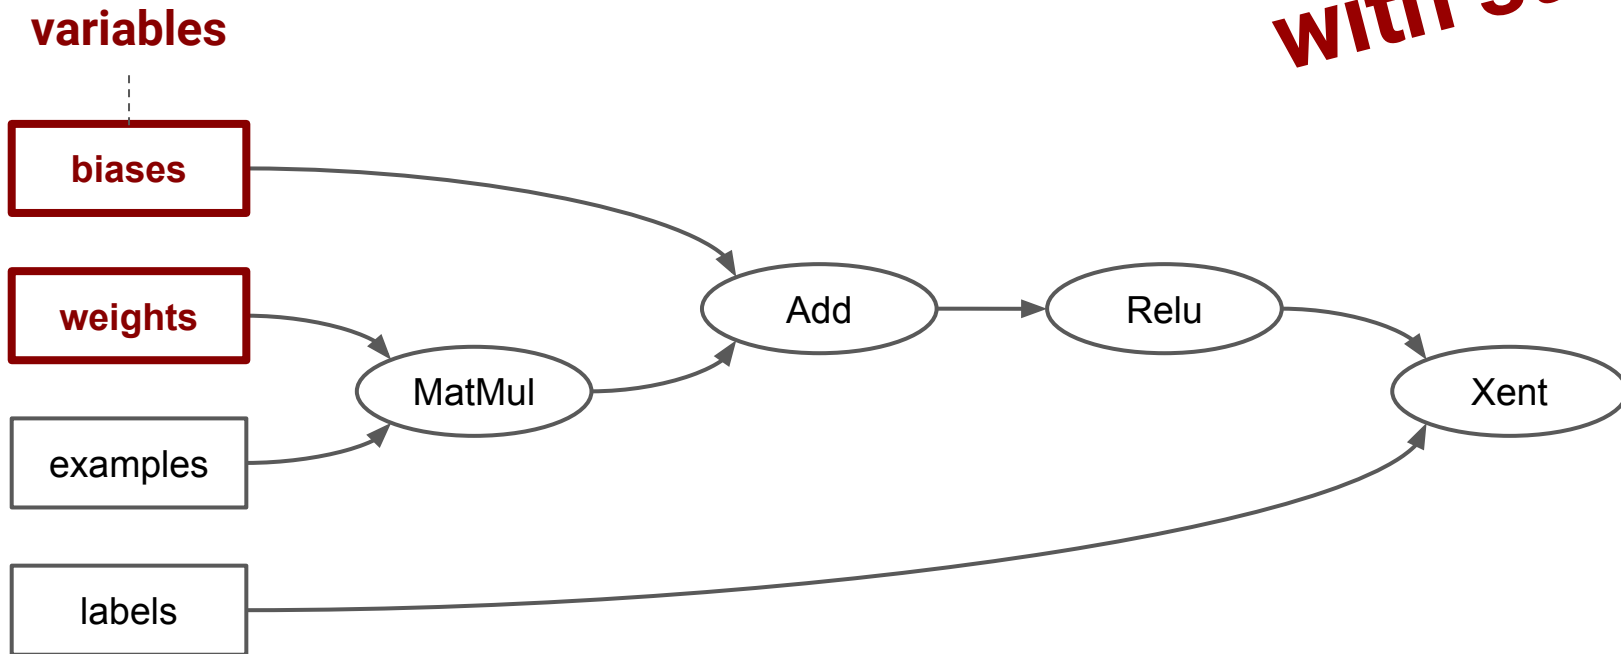


Any Computation is a TensorFlow Graph



Any Computation is a TensorFlow Graph


with state



Any Computation is a TensorFlow Graph

Simple gradient descent:

with state



Linear Regression


The background of the slide features a solid orange gradient. In the lower half, there is a stylized representation of a mountain range using a white wireframe mesh. The mesh lines form a grid that follows the contours of the peaks and valleys, creating a 3D effect against the orange backdrop.

Linear Regression

Diagram illustrating the linear regression equation $y = Wx + b$.

The equation components are labeled with arrows:

- result** points to y .
- input** points to x .
- parameters** points to W and b .



What are we trying to do?

Mystery equation: $y = 0.1 * x + 0.3 + \text{noise}$

Model: $y = W * x + b$

Objective: Given enough (x, y) value samples, figure out the value of W and b .

$y = Wx + b$ in TensorFlow



```
import tensorflow as tf
```

A TensorFlow graph consists of the following parts which will be detailed below:

1.Placeholder variables used to change the input to the graph.

2.Model variables that are going to be optimized so as to make the model perform better.

3.The model which is essentially just a mathematical function that calculates some output given the input in the placeholder variables and the model variables.

4.A cost measure that can be used to guide the optimization of the variables.

5.An optimization method which updates the variables of the model.

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

```
x = tf.placeholder(shape=[None],  
                   dtype=tf.float32, name="x")
```

Placeholder variables serve as the input to the graph that we may change each time we execute the graph.

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```

```
x = tf.placeholder(shape=[None],  
                   dtype=tf.float32, name="x")
```

```
W = tf.get_variable(shape=[], name="W")
```

Apart from the placeholder variables that were defined above and which serve as feeding input data into the model, there are also some model variables that must be changed by TensorFlow so as to make the model perform better on the training data.

$y = Wx + b$ in TensorFlow

```
import tensorflow as tf
```



```
x = tf.placeholder(shape=[None],  
                    dtype=tf.float32, name="x")
```

Text

```
W = tf.get_variable(shape=[], name="W")
```

```
b = tf.get_variable(shape=[], name="b")
```

$y = Wx + b$ in TensorFlow


```
import tensorflow as tf
```

```
x = tf.placeholder(shape=[None],  
                    dtype=tf.float32, name="x")
```

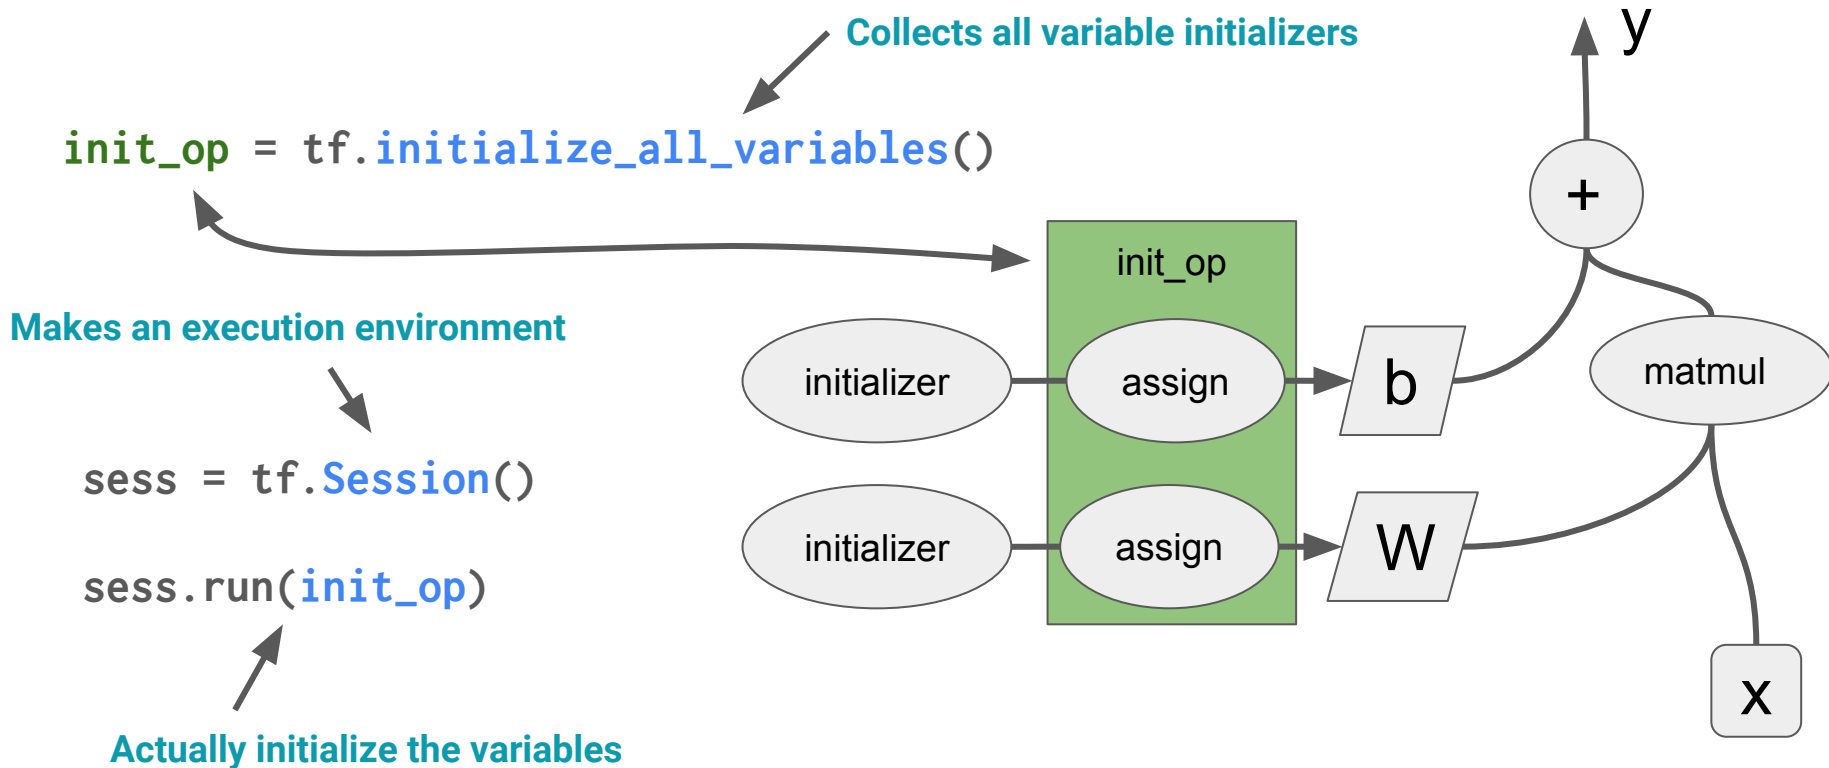
```
W = tf.get_variable(shape=[], name="W")
```

```
b = tf.get_variable(shape=[], name="b")
```

```
y = W * x + b
```



Variables Must be Initialized




Running the Computation

`x_in = 3`

`sess.run(y, feed_dict={x: x_in})`

- Only what's used to compute a fetch will be evaluated
- All Tensors can be fed, but all placeholders must be fed



Putting it all together

```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                   dtype=tf.float32,
                   name='x')

W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b
```

```
with tf.Session() as sess:
```

```
    sess.run(tf.initialize_all_variables())
```

```
    print(sess.run(y, feed_dict={x: x_in}))
```

Build the graph

Prepare execution environment

Initialize variables

Run the computation (usually often)


Define a Loss

Given x , y compute a loss, for instance:

$$L = (y - y_{label})^2$$

```
# create an operation that calculates loss.  
loss = tf.reduce_mean(tf.square(y - y_data))
```

Minimize loss: optimizers



`tf.train.AdadeltaOptimizer`

`tf.train.AdagradOptimizer`

`tf.train.AdagradDAOptimizer`

`tf.train.AdamOptimizer`


...

Train

Feed (x, y_{label}) pairs and adjust W and b to decrease the loss.

$$W \leftarrow W - \eta (dL/dW)$$

$$b \leftarrow b - \eta (dL/db)$$



TensorFlow computes
gradients automatically

```
# Create an optimizer
```

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

```
# Create an operation that minimizes loss.
```



Learning rate

```
train = optimizer.minimize(loss)
```

Putting it all together

```
loss = tf.reduce_mean(tf.square(y - y_label))
```

} Define a loss

```
optimizer = tf.train.GradientDescentOptimizer(0.5)
```

} Create an optimizer

```
train = optimizer.minimize(loss)
```

} Op to minimize the loss

```
with tf.Session() as sess:
```

```
    sess.run(tf.initialize_all_variables())
```

} Initialize variables


```
    for i in range(1000):
```

```
        sess.run(train, feed_dict={x: x_in[i],
```

```
                                   y_label: y_in[i]}))
```

} Iteratively run the training op

TensorBoard



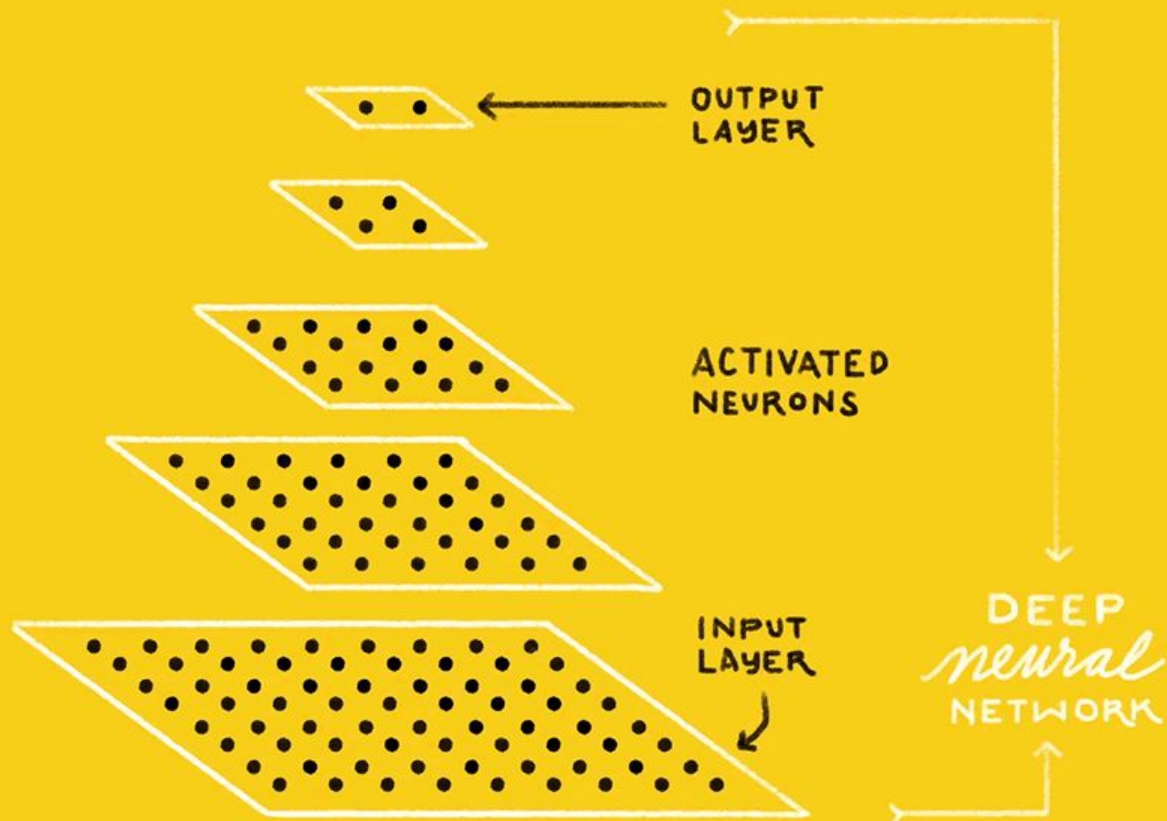
Deep Neural Network

The background features a solid orange gradient. Overlaid on this is a white wireframe mesh that forms a series of undulating, mountain-like peaks and valleys, creating a 3D effect.

IS THIS A
CAT or DOG?

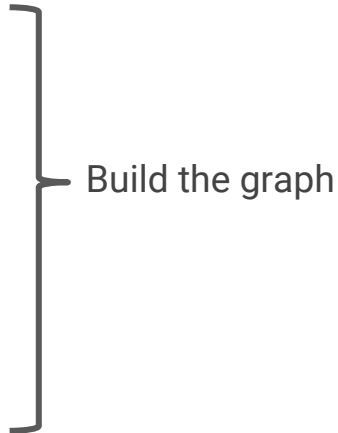


CAT DOG



Remember linear regression?


```
import tensorflow as tf
x = tf.placeholder(shape=[None],
                   dtype=tf.float32,
                   name='x')
W = tf.get_variable(shape=[], name='W')
b = tf.get_variable(shape=[], name='b')
y = W * x + b
loss = tf.reduce_mean(tf.square(y - y_label))
optimizer = tf.train.GradientDescentOptimizer(0.5)
train = optimizer.minimize(loss)
...
```




Build the graph

Convolutional DNN

```
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.conv2d(x, kernel_size=[5,5], ...)
x = tf.contrib.layers.max_pool2d(x, kernel_size=[2,2], ...)
x = tf.contrib.layers.fully_connected(x, activation_fn=tf.nn.relu)
x = tf.contrib.layers.dropout(x, 0.5)
logits = tf.config.layers.linear(x)
```



Defining Complex Networks



Tutorials & Courses

Tutorials on [tensorflow.org](https://www.tensorflow.org):

Image recognition: https://www.tensorflow.org/tutorials/image_recognition

Word embeddings: <https://www.tensorflow.org/versions/word2vec>

Language Modeling: <https://www.tensorflow.org/tutorials/recurrent>

Translation: <https://www.tensorflow.org/versions/seq2seq>

Deep Dream:

<https://tensorflow.org/code/tensorflow/examples/tutorials/deepdream/deepdream.ipynb>

The background is a solid orange color. In the lower half, there is a white wireframe illustration of a mountain range. The mountains are composed of a grid of lines that form a series of peaks and valleys, creating a three-dimensional effect. The word "Extras" is written in a bold, white, sans-serif font on the left side of the image, partially overlapping the wireframe mountains.

Extras


Inception



An Alaskan Malamute (left) and a Siberian Husky (right). Images from Wikipedia.

<https://research.googleblog.com/2016/08/improving-inception-and-image.html>


Show and Tell



A person on a beach flying a kite.

<https://research.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>

Parsey McParseface



<https://research.googleblog.com/2016/05/announcing-syntaxnet-worlds-most.html>

Text Summarization

Original text

- *Alice and Bob took the train to visit the zoo. They saw a **baby giraffe, a lion, and a flock of colorful tropical birds.***

Abstractive summary

- *Alice and Bob visited the zoo and saw **animals and birds.***

Mobile TensorFlow

TensorFlow was designed with mobile and embedded platforms in mind. We have sample code and build support you can try now for these platforms:

[Android](#)

[iOS](#)

[Raspberry Pi](#)

Many applications can benefit from on-device processing. Google Translate's instant visual translation is a great example. By running its processing locally, users get an incredibly responsive and interactive experience.

Mobile TensorFlow makes sense when there is a poor or missing network connection, or where sending continuous data to a server would be too expensive. We are working to help developers make lean mobile apps using TensorFlow, both by continuing to reduce the code footprint, and supporting [quantization](#) and [lower precision arithmetic](#) that reduce model size.



Monet - Sunflowers: 0.992



